Băhēm

Provably Secure Symmetric Cipher

M. Rajululkahf ©¹
May 2, 2022

Overview

This paper proposes Băhēm; a symmetric cipher that, when given a random-looking key \mathbf{k} , a true random number generator (TRNG) and a cleartext message \mathbf{m} to encrypt, no cryptanalysis can degrade its security below $\min(H(\mathbf{m}), H(\mathbf{k}))$ bits of entropy, even under Grover's algorithm [1] or even if it turned out that P = NP.

Aside from the cost of memory access and input/output processing, Băhēm is also highly parallelise-able, and requires only three additions and one bitwise exclusive-or operation (XOR) in order to encrypt or decrypt.

Its early prototype, Alyal, achieved similar runtime speeds to OpenSSL's ChaCha20 [2]; slightly faster decryption, while slightly slower encryption when the TRNG was prepared in a file in advance. Future versions, with better TRNG optimisations, should be able to enable the prototype to have faster run-time for both, encryption and decryption, alike. Either way, Alyal is highly practical, specially when considering its strong security guarantee. Further speed gain is possible when Alyal utilises Băhēm's high parallelism.

Notation

 $H(\mathbf{x})$: Shannon's entropy of random variable \mathbf{x} .

 $\mathbf{x} + \mathbf{y} \mod 2^{128}$: Unsigned 128-bit addition.

random(128): A sequence of 128 many random bits generated by a TRNG.

k: A 128-bit pre-shared secret key with enough H(k) that looks random. Ideally k = random(128).

 \mathbf{m} : An arbitrarily-long cleartext message of $|\mathbf{m}|$ many bits.

 $\lceil \frac{|\mathbf{m}|}{128} \rceil$: Number of 128-bit blocks in clear text \mathbf{m} .

 \mathbf{m}_b : The b^{th} 128-bit block from \mathbf{m} .

 $\mathbf{p}_b = \operatorname{random}(128), \mathbf{q}_b = \operatorname{random}(128)$: A pair of uniformly distributed 128 many random bits.

 $\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b$: Encrypted forms of $\mathbf{p}_b, \mathbf{q}_b$ and \mathbf{m}_b , respectively.

Contents

1	Proposed Algorithm Security Analysis									
2										
3	Implementation Examples									
	3.1	.1 C Functions								
	3.2	A File	Encryption Tool	3						
		3.2.1	Installation	3						
		3.2.2	Usage	3						
		3.2.3	Benchmark	3						
4	Cor	nclusion	18	4						

1 Proposed Algorithm

Algorithms 1 and 2 show Băhēm's encryption and decryption by which the process is repeated over every 128-bit blocks of \mathbf{m} .

Algorithm 1: Băhēm encryption

input: $k, m_0, m_1, ...$ output: $(\hat{\mathbf{p}}_0, \hat{\mathbf{q}}_0, \hat{\mathbf{m}}_0), (\hat{\mathbf{p}}_1, \hat{\mathbf{q}}_1, \hat{\mathbf{m}}_1), ...$

for $b \in (0, 1, ..., \lceil \frac{|\mathbf{m}|}{128} \rceil - 1)$ do $\begin{vmatrix} \mathbf{p}_b \leftarrow \text{random}(128) \\ \mathbf{q}_b \leftarrow \text{random}(128) \\ \hat{\mathbf{p}}_b \leftarrow \mathbf{p}_b + \mathbf{k} \mod 2^{128} \\ \hat{\mathbf{q}}_b \leftarrow \mathbf{q}_b + \mathbf{k} \mod 2^{128} \\ \hat{\mathbf{m}}_b \leftarrow \mathbf{m}_b \oplus (\mathbf{p}_b + \mathbf{q}_b \mod 2^{128}) \end{vmatrix}$

Algorithm 2: Băhēm decryption

input : k, $(\hat{p}_0, \hat{q}_0, \hat{m}_0)$, $(\hat{p}_1, \hat{q}_1, \hat{m}_1)$, ... output: $m_0, m_1, ...$

for $b \in (0, 1, ..., \lceil \frac{|\mathbf{m}|}{128} \rceil - 1)$ do $\begin{array}{c} \mathbf{p}_b \leftarrow \hat{\mathbf{p}}_b - \mathbf{k} \bmod 2^{128} \\ \mathbf{q}_b \leftarrow \hat{\mathbf{q}}_b - \mathbf{k} \bmod 2^{128} \\ \mathbf{m}_b \leftarrow \hat{\mathbf{m}}_b \oplus (\mathbf{p}_b + \mathbf{q}_b \bmod 2^{128}) \end{array}$

This work is licensed under a Creative Commons "Attribution 4.0 International" license.



¹Author's e-mail address: {last name}@pm.me

2 Security Analysis

The Băhēm encryption is essentially the XOR cryptosystem:

$$\hat{\mathbf{m}}_b \leftarrow \mathbf{m}_b \oplus \underbrace{\left(\mathbf{p}_b + \mathbf{q}_b \bmod 2^{128}\right)}_{\text{Encryption pad}}$$

It trivially follows from Shannon's perfect secrecy proof of the one-time pad (OTP) [3] that Băhēm is secure if and only if its encryption pad maintains its key's entropy as shown in eq. (1), even if the adversary knows $\hat{\mathbf{p}}_b$, $\hat{\mathbf{q}}_b$, $\hat{\mathbf{m}}_b$ and the cleartext message \mathbf{m}_b .

$$H(\mathbf{k}|\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b, \mathbf{m}_b)$$

$$= H(\mathbf{k}|\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \mathbf{p}_b + \mathbf{q}_b \mod 2^{128})$$

$$= H(\mathbf{k})$$
(1)

To simplify the analysis, suppose that the size of a block in Băhēm is 3 bits only, and that the cleartext block \mathbf{m}_b is known to the adversary, which implies that the adversary can trivially know that:

$$\mathbf{p}_b + \mathbf{q}_b \mod 2^3 = \hat{\mathbf{m}}_b \oplus \mathbf{m}_b$$

in addition to adversary's knowledge of the public variables $\hat{\mathbf{p}}_b$ and $\hat{\mathbf{q}}_b$. More specifically, suppose that the adversary found that:

$$0 = \hat{\mathbf{p}}_b = \mathbf{p}_b + \mathbf{k} \mod 2^3$$
$$3 = \hat{\mathbf{q}}_b = \mathbf{q}_b + \mathbf{k} \mod 2^3$$
$$5 = \mathbf{p}_b + \mathbf{q}_b \mod 2^3$$

Then, the question is: will this information reduce the space from which the key \mathbf{k} is chosen from? In other words, what are the possible values of \mathbf{k} that can lead to the outputs 0, 3 and 5 above? Table 1 visualises this.

As shown in table 1, the total number of horizontal, or vertical, intersections that simultaneously cross all of the outputs 0, 3 and 5, remain 2^3 . Meaning, the total number of values of \mathbf{k} that could lead to the outputs remains 2^3 .

This 3-bit example can be trivially extended by induction to show that the same conclusions hold even with a 128-bit unsigned addition and any other output numbers than 0, 3 and 5.

Therefore, we can conclude that adversary's knowledge of the public variables $\hat{\mathbf{p}}_b$, $\hat{\mathbf{q}}_b$, $\hat{\mathbf{m}}_b$ and the clear-text \mathbf{m}_b , which leads to deducing $\mathbf{p}_b + \mathbf{q}_b \mod 2^{128}$, can not reduce $\mathbf{H}(\mathbf{k})$.

Since $\hat{\mathbf{p}}_b$, $\hat{\mathbf{q}}_b$, $\hat{\mathbf{m}}_b$ and the clear text \mathbf{m}_b are exhaustively all of the outputs of Băhēm that can be accessible to an adversary, and since they do not reduce

				\mathcal{Y}				
	0	1	2	3	4	5	6	7
0	0	1	2	3	4	5	6	7
1	1	2	3	4	5	6	7	0
2	2	3	4	5	6	7	0	1
\mathcal{X} 3	3	4	5	6	7	0	1	2
4	4	5	6	7	0	1	2	3
5	5	6	7	0	1	2	3	4
6	6	7	0	1	2	3	4	5
7	7	0	1	2	3	4	5	6

Table 1: Exhaustive unsigned 3-bit addition. For a given output $\mathbf{x}+\mathbf{y} \mod 2^3$, there are 2^3 many possible input values of $(\mathbf{x},\mathbf{y}) \in \mathcal{X} \times \mathcal{Y}$ that map to $\mathbf{x}+\mathbf{y} \mod 2^3$

Băhēm's key space, therefore no cryptanalysis can reduce the entropy Băhēm's encryption pad below $H(\mathbf{k})$.

Lemma 1 (Secure private values).

$$H(\mathbf{k}|\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b, \mathbf{m}_b)$$

$$= H(\mathbf{p}|\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b, \mathbf{m}_b)$$

$$= H(\mathbf{q}|\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b, \mathbf{m}_b)$$

$$= H(\mathbf{k})$$

Lemma 2 (Secure encryption pad).

$$H(\mathbf{p} + \mathbf{q} \mod 2^{128} | \hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b, \mathbf{m}_b) = H(\mathbf{k})$$

Since Băhēm is an XOR cryptosystem, and since its encryption pad is secure (lemma 2), therefore it has to follow by Shannon's perfect secrecy [3] that Băhēm's encryption is secure as well.

Theorem 1 (Secure encryption).

$$H(\mathbf{m}_b|\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b) = \min(H(\mathbf{m}_b), H(\mathbf{k}))$$

Note 1. Băhēm does not have perfect secrecy, as this requires an impractical key that is as long as the message itself. Băhēm rather has a $H(\mathbf{k})$ -bit security. Shannon's proof of perfect secrecy of the OTP is cited only for its relevance in proving Băhēm's $H(\mathbf{k})$ -bit security as both of them are XOR cryptosystems.

3 Implementation Examples

3.1 C Functions

In this example, the caller is expected to initialise a 128 bits key \mathbf{k} , a pair of random pads, \mathbf{p} and \mathbf{q} , each

of which is $\texttt{len} \times 128$ bits long, in order to encrypt a $\texttt{len} \times 64$ bits long cleartext message \mathbf{m} , where len is an even number. The encryption happens in-place, so the caller does not have to allocate separate memory for the output. Since 128-bit wide CPU instructions are not common, this code operates in 64-bit basis, each time with a different 64-bit part of the key.

```
void baheem_enc(
   uint64_t *k, /* 128bit pre-shared key */
   uint64_t *p, /* random pad 1
   uint64_t *q, /* random pad 2
    uint64_t *m, /* message
    size_t len /* length of m = p = q
) {
    size_t i;
    for (i = 0; i < len; i += 2) {
               ^= p[i]
        m[i]
                         + q[i];
        m[i+1] = p[i+1] + q[i+1];
        p[i]
               += k[0];
        q[i]
               += k[0];
        p[i+1] += k[1];
        q[i+1] += k[1];
    }
}
```

Likewise, the following is an example implementing the corresponding in-place decryption function.

3.2 A File Encryption Tool

Alyal is an single-threaded implementation to demonstrate Băhēm's practical utility with real-world scenarios. Internally, Alyal uses the baheem_enc and baheem_dec functions that were presented earlier in this section.

3.2.1 Installation

```
git clone \
  https://codeberg.org/rajululkahf/alyal
cd alyal
make
make test
```

3.2.2 Usage

```
alyal (enc|dec) IN OUT [TRNG]
alyal help
```

To encrypt a cleartext file **a** and save it as file **b**:

```
alyal enc a b
```

To decrypt the latter back to its cleartext form and save it as file c:

```
alyal dec b c
```

3.2.3 Benchmark

This is a benchmark that was performed on a computer with a 3.4GHz Intel Core i5-3570K CPU, 32GB RAM, 7200 RPM hard disks, Linux 5.17.4-gentoo-x86-64, and OpenSSL 1.1.1n.

	OpenSSL	Alyal		
	ChaCha20	${ m Breve{a}har{e}m}$		
		/dev/random	file.rand	
Encrypt	$0.87 \mathrm{\ secs}$	$3.91 \mathrm{\ secs}$	1.40 secs	
500MB	1.04 secs	4.25 secs	1.82 secs	
	1.04 secs	$4.27 \mathrm{secs}$	1.73 secs	
Decrypt	$0.90 \mathrm{secs}$	$0.64~{ m secs}$		
500MB	1.06 secs	$0.89 \mathrm{secs}$		
	1.06 secs	0.81	secs	

Table 2: Wall-clock run-time comparison between OpenSSL's ChaCha20, and Alyal's Băhēm implementation with two sources as the TRNG: /dev/random and file.rand; the latter is simply /dev/random that was prepared in advance.

Table 2 shows that, while the early Băhēm prototype, Alyal, has a faster decryption run-time than OpenSSL's ChaCha20, it has slower encryption run-time. However:

- The differences in run-time are insignificant for most applications, which proves Băhēm's practical utility in the real world.
- 2. Băhēm's provable security should arguably justify waiting the 3 extra seconds for the 500MB data, specially that many user applications involve encrypting much smaller data sizes with unnoticeable time difference
- 3. Preparing the random bits in advance significantly reduces the encryption time as shown with the file.rand case in table 2, and can be optimised further should it be prepared in memory.

4. Alyal is currently single-threaded despite Băhēm's capacity for high parallelism as all blocks are independent. This gives room for future versions to be significantly faster.

4 Conclusions

This paper proposed Băhēm with the following properties:

Secure. No cryptanalysis can degrade its security below $min(H(\mathbf{m}), H(\mathbf{k}))$ bits.

Fast. Requires only three additions and a single XOR per encryption or decryption alike.

Highly parallelisable as the encryption, or decryption, of any bit is independent of other bits.

Băhēm's single-threaded prototype, Alyal, outperformed OpenSSL's ChaCha20 when decrypting files, despite Băhēm's 2 bits overhead, which demonstrates that such overhead is negligible in practice.

While the prototype has a slower encryption runtime due to its use of a TRNG, optimising it is trivial by preparing the TRNG in advance.

Simple. Băhēm's simplicity implies fewer expected number of implementation bugs, and therefore higher practical security.

References

- [1] Lov K. Grover. A fast quantum mechanical algorithm for database search. In *Proceedings of the Twenty-Eighth Annual ACM Symposium on Theory of Computing*, STOC '96, page 212–219, New York, NY, USA, 1996. Association for Computing Machinery.
- [2] Daniel Bernstein. ChaCha, a variant of Salsa20. 01 2008.
- [3] C. E. Shannon. Communication theory of secrecy systems. *The Bell System Technical Journal*, 28(4):656–715, 1949.