

Bähēm

Provably Secure Symmetric Cipher

M. Rajululkahf¹

April 29, 2022

Overview

This paper proposes Bähēm; a symmetric cipher that, when given a random-looking key \mathbf{k} , a true random number generator (TRNG) and a cleartext message \mathbf{m} to encrypt, no cryptanalysis can degrade its security below $\min(H(\mathbf{m}), H(\mathbf{k}))$ bits of entropy, even under Grover's algorithm [1] or even if it turned out that $P = NP$.

Bähēm is also highly parallelise-able, and requires only three additions and a single bitwise exclusive-or operation (XOR).

Its early prototype, Alyal, achieved similar run-time speeds to OpenSSL's ChaCha20 [2]; slightly faster decryption, while slightly slower encryption when the TRNG was prepared in a file in advance. Future versions, with better TRNG optimisations, should be able to enable the prototype to have faster run-time for both, encryption and decryption, alike.

Notation

$H(\mathbf{x})$: Shannon's entropy of random variable \mathbf{x} .

$\mathbf{x} + \mathbf{y} \bmod 2^{128}$: Unsigned 128-bit addition.

$\text{random}(128)$: A sequence of 128 many random bits generated by a TRNG.

\mathbf{k} : A 128-bit pre-shared secret key with enough $H(\mathbf{k})$ that looks random. Ideally $\mathbf{k} = \text{random}(128)$.

\mathbf{m} : An arbitrarily-long cleartext message of $|\mathbf{m}|$ many bits.

$\lceil \frac{|\mathbf{m}|}{128} \rceil$: Number of 128-bit blocks in cleartext \mathbf{m} .

\mathbf{m}_b : The b^{th} 128-bit block from \mathbf{m} .

$\mathbf{p}_b = \text{random}(128)$, $\mathbf{q}_b = \text{random}(128)$: A pair of uniformly distributed random bits. This is generated dynamically by the implementation, transparently from the user, for every block.

$\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b$: Encrypted forms of \mathbf{p}_b , \mathbf{q}_b and \mathbf{m}_b , respectively.

Contents

1	Proposed Algorithm	1
2	Security Analysis	2
3	Implementation Examples	2
3.1	C Functions	2
3.2	Alyal	3
3.2.1	Installation	3
3.2.2	Benchmark	3
4	Conclusions	3

1 Proposed Algorithm

Algorithms 1 and 2 show Bähēm's encryption and decryption by which the process is repeated over every 128-bit blocks of \mathbf{m} .

Algorithm 1: Bähēm encryption

input : $\mathbf{k}, \mathbf{m}_0, \mathbf{m}_1, \dots$

output: $(\hat{\mathbf{p}}_0, \hat{\mathbf{q}}_0, \hat{\mathbf{m}}_0), (\hat{\mathbf{p}}_1, \hat{\mathbf{q}}_1, \hat{\mathbf{m}}_1), \dots$

```
for  $b \in (0, 1, \dots, \lceil \frac{|\mathbf{m}|}{128} \rceil - 1)$  do
   $\mathbf{p}_b \leftarrow \text{random}(128)$ 
   $\mathbf{q}_b \leftarrow \text{random}(128)$ 
   $\hat{\mathbf{p}}_b \leftarrow \mathbf{p}_b + \mathbf{k} \bmod 2^{128}$ 
   $\hat{\mathbf{q}}_b \leftarrow \mathbf{q}_b + \mathbf{k} \bmod 2^{128}$ 
   $\hat{\mathbf{m}}_b \leftarrow \mathbf{m}_b \oplus (\mathbf{p}_b + \mathbf{q}_b \bmod 2^{128})$ 
```

Algorithm 2: Bähēm decryption

input : $\mathbf{k}, (\hat{\mathbf{p}}_0, \hat{\mathbf{q}}_0, \hat{\mathbf{m}}_0), (\hat{\mathbf{p}}_1, \hat{\mathbf{q}}_1, \hat{\mathbf{m}}_1), \dots$

output: $\mathbf{m}_0, \mathbf{m}_1, \dots$

```
for  $b \in (0, 1, \dots, \lceil \frac{|\mathbf{m}|}{128} \rceil - 1)$  do
   $\mathbf{p}_b \leftarrow \hat{\mathbf{p}}_b - \mathbf{k} \bmod 2^{128}$ 
   $\mathbf{q}_b \leftarrow \hat{\mathbf{q}}_b - \mathbf{k} \bmod 2^{128}$ 
   $\mathbf{m}_b \leftarrow \hat{\mathbf{m}}_b \oplus (\mathbf{p}_b + \mathbf{q}_b \bmod 2^{128})$ 
```

¹Author's e-mail address: {last name}@pm.me



2 Security Analysis

The Bähēm encryption is essentially the XOR cryptosystem:

$$\hat{\mathbf{m}}_b \leftarrow \mathbf{m}_b \oplus \underbrace{(\mathbf{p}_b + \mathbf{q}_b \bmod 2^{128})}_{\text{Encryption pad}}$$

It follows from Shannon’s perfect secrecy proof of the one-time pad (OTP) [3] that Bähēm is secure if and only if its encryption pad maintains its key’s entropy as shown in eq. (1), even if the adversary knows $\hat{\mathbf{p}}_b$, $\hat{\mathbf{q}}_b$, $\hat{\mathbf{m}}_b$ and the cleartext message \mathbf{m}_b .

$$\begin{aligned} H(\mathbf{k}|\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b, \mathbf{m}_b) \\ &= H(\mathbf{k}|\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \mathbf{p}_b + \mathbf{q}_b \bmod 2^{128}) \\ &= H(\mathbf{k}) \end{aligned} \quad (1)$$

To simplify the analysis, suppose that the size of a block in Bähēm is 3 bits only, and that the cleartext block \mathbf{m}_b is known to the adversary, which implies that the adversary can trivially know that:

$$\mathbf{p}_b + \mathbf{q}_b \bmod 2^3 = \hat{\mathbf{m}}_b \oplus \mathbf{m}_b$$

in addition to adversary’s knowledge of the public variables $\hat{\mathbf{p}}_b$ and $\hat{\mathbf{q}}_b$. More specifically, suppose that the adversary found that:

$$\begin{aligned} 0 &= \hat{\mathbf{p}}_b = \mathbf{p}_b + \mathbf{k} \bmod 2^3 \\ 3 &= \hat{\mathbf{q}}_b = \mathbf{q}_b + \mathbf{k} \bmod 2^3 \\ 5 &= \mathbf{p}_b + \mathbf{q}_b \bmod 2^3 \end{aligned}$$

Then, the question is: will this information reduce the space from which the key \mathbf{k} is chosen from? In other words, what are the possible values of \mathbf{k} that can lead to the outputs 0, 3 and 5 above? Table 1 visualises this.

As shown in table 1, the total number of horizontal and vertical intersections that simultaneously cross all of the outputs 0, 3 and 5, remain 2^3 . Meaning, the total number of values of \mathbf{k} that could lead to the outputs remains 2^3 .

This 3-bit example can be trivially extended by induction to show that the same conclusions hold even with a 128-bit unsigned addition and any other output numbers than 0, 3 and 5.

Therefore, we can conclude that adversary’s knowledge of the public variables $\hat{\mathbf{p}}_b$, $\hat{\mathbf{q}}_b$, $\hat{\mathbf{m}}_b$ and the cleartext \mathbf{m}_b , which leads to deducing $\mathbf{p}_b + \mathbf{q}_b \bmod 2^{128}$, can not reduce $H(\mathbf{k})$.

Since $\hat{\mathbf{p}}_b$, $\hat{\mathbf{q}}_b$, $\hat{\mathbf{m}}_b$ and the cleartext \mathbf{m}_b are exhaustively all of the outputs of Bähēm that can be accessible to an adversary, and since they do not reduce

	\mathcal{Y}							
	0	1	2	3	4	5	6	7
0	0	1	2	3	4	5	6	7
1	1	2	3	4	5	6	7	0
2	2	3	4	5	6	7	0	1
3	3	4	5	6	7	0	1	2
4	4	5	6	7	0	1	2	3
5	5	6	7	0	1	2	3	4
6	6	7	0	1	2	3	4	5
7	7	0	1	2	3	4	5	6

Table 1: Exhaustive unsigned 3-bit addition. For a given output $\mathbf{x} + \mathbf{y} \bmod 2^3$, there are 2^3 many possible input values of $(\mathbf{x}, \mathbf{y}) \in \mathcal{X} \times \mathcal{Y}$ that map to $\mathbf{x} + \mathbf{y} \bmod 2^3$.

Bähēm’s key space, therefore no cryptanalysis can reduce Bähēm’s security below $H(\mathbf{k})$.

Theorem 2.1 (Secure encryption pad). *If the key \mathbf{k} looks like a uniformly distributed 128-bit random number, and \mathbf{p}_b and \mathbf{q}_b are 128-bit random numbers generated by a TRNG, then Bähēm guarantees that its encryption pad satisfies:*

$$H(\mathbf{k}|\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b, \mathbf{m}_b) = H(\mathbf{k})$$

Since Bähēm is an XOR cryptosystem, and since its encryption pad is secure (theorem 2.1), therefore it has to follow by Shannon’s perfect secrecy [3] that Bähēm’s encryption is secure as well.

Theorem 2.2 (Secure encryption).

$$H(\mathbf{m}_b|\hat{\mathbf{p}}_b, \hat{\mathbf{q}}_b, \hat{\mathbf{m}}_b) = \min(H(\mathbf{m}_b), H(\mathbf{k}))$$

Note 2.1. *This paper does not claim that Bähēm has perfect secrecy, but rather claims that it has a $H(\mathbf{k})$ -bit security. Shannon’s perfect secrecy proof is cited only for its relevance in proving Bähēm’s $H(\mathbf{k})$ -bit security as the OTP is nonetheless a similar XOR cryptosystem.*

3 Implementation Examples

3.1 C Functions

In this example, the caller is expected to initialise a 128 bits key \mathbf{k} , a pair of random pads, \mathbf{p} and \mathbf{q} , each of which is $1\text{en} \times 64$ bits long, in order to encrypt a $1\text{en} \times 64$ bits long cleartext message \mathbf{m} . The encryption happens in-place, so the caller does not have to allocate separate memory for the ciphertext.

```

void baheem_enc(
    uint64_t *k, /* 128bit pre-shared key */
    uint64_t *p, /* random pad 1 */
    uint64_t *q, /* random pad 2 */
    uint64_t *m, /* message */
    size_t len /* length of m = p = q */
) {
    size_t i;
    for (i = 0; i < len; i++) {
        m[i] ^= p[i] + q[i];
        p[i] += k[0];
        q[i] += k[1];
    }
}

```

Likewise, the following is an example implementing the corresponding in-place decryption function.

```

void baheem_dec(... same input ...) {
    size_t i;
    for (i = 0; i < len; i++) {
        p[i] -= k[0];
        q[i] -= k[1];
        m[i] ^= p[i] + q[i];
    }
}

```

3.2 Alyal

Alyal is an early single-threaded implementation to demonstrate Bähēm’s practical utility with real-world scenarios. Internally, Alyal uses the `baheem_enc` and `baheem_dec` functions that were presented earlier in this section.

3.2.1 Installation

```

> git clone \
    https://codeberg.org/rajululkahf/alyal
> cd alyal
> make
> dd bs=1MB count=500 \
    if=/dev/zero of=test.txt
> ./alyal enc test.txt test.enc
> ./alyal dec test.enc test.enc.txt
> shasum *

```

3.2.2 Benchmark

This benchmark that was performed on a machine with a 3.4GHz Intel Core i5-3570K CPU, 32GB RAM, 7200 RPM hard disks, Linux 5.17.4-gentoo-x86-64, and OpenSSL 1.1.1n.

Table 2 shows that, while the early Bähēm prototype, Alyal, has a faster decryption run-time than

	OpenSSL ChaCha20	Alyal Bähēm	
		/dev/random	file.rand
Encrypt	0.87 secs	3.91 secs	1.40 secs
500MB	1.04 secs	4.25 secs	1.82 secs
	1.04 secs	4.27 secs	1.73 secs
Decrypt	0.90 secs	0.64 secs	
500MB	1.06 secs	0.89 secs	
	1.06 secs	0.81 secs	

Table 2: Wall-clock run-time comparison between OpenSSL’s ChaCha20, and Alyal’s Bähēm implementation with two sources as the TRNG: `/dev/random` and `file.rand`; the latter is simply `/dev/random` that was prepared in advance.

OpenSSL’s ChaCha20, it has slower encryption run-time. However:

1. The differences in run-time are insignificant for most applications, which proves Bähēm’s practical utility in the real world.
2. Bähēm’s provable security should arguably justify waiting the 3 extra seconds for the 500MB data, specially that many user applications involve encrypting much smaller data sizes with unnoticeable time difference
3. Preparing the random bits in advance significantly reduces the delay as shown with the `file.rand` case, and can be optimised further should it be prepared in memory.
4. Alyal is currently single-threaded despite Bähēm’s capacity for high parallelism as all blocks are independent. This gives room for future versions to be significantly faster.

4 Conclusions

This paper proposed Bähēm with the following properties:

Secure. No cryptanalysis can degrade its security below $\min(H(\mathbf{m}), H(\mathbf{k}))$ bits.

Fast. Requires only three additions and a single XOR per encryption or decryption alike.

Highly parallelisable as the encryption, or decryption, of any bit is independent of other bits.

Bähēm’s single-threaded prototype (Alyal) outperformed OpenSSL’s ChaCha20 when decrypting files, despite Bähēm’s 2 bits overhead, which

demonstrates that such overhead is negligible in practice.

While the prototype has a slower encryption runtime due to its use of a TRNG, optimising it is trivial by preparing the TRNG in advance.

Simple. Bähēm’s simplicity implies fewer expected number of implementation bugs, and therefore higher practical security.

References

- [1] Lov K. Grover. A fast quantum mechanical algorithm for database search. In *Proceedings of the Twenty-Eighth Annual ACM Symposium on Theory of Computing, STOC '96*, page 212–219, New York, NY, USA, 1996. Association for Computing Machinery.
- [2] Daniel Bernstein. Chacha, a variant of salsa20. 01 2008.
- [3] C. E. Shannon. Communication theory of secrecy systems. *The Bell System Technical Journal*, 28(4):656–715, 1949.